

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import AdaBoostClassifier

```

```
data = pd.read_csv('Mall_Customers.csv')
```

```
data
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	
39					
1	2	Male	21	15	
81					
2	3	Female	20	16	
6					
3	4	Female	23	16	
77					
4	5	Female	31	17	
40					
..	...	...	...	...	..
195	196	Female	35	120	
79					
196	197	Female	45	126	
28					
197	198	Male	32	126	
74					
198	199	Male	32	137	
18					
199	200	Male	30	137	
83					

```
[200 rows x 5 columns]
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 200 entries, 0 to 199
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	CustomerID	200 non-null	int64
1	Gender	200 non-null	object
2	Age	200 non-null	int64

```
3   Annual Income (k$)      200 non-null    int64
4   Spending Score (1-100)  200 non-null    int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
data.describe()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
data.isnull().sum()
```

```
CustomerID      0
Gender           0
Age             0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64
```


```
sns.countplot(data.Gender)
```

```
D:\Anaconda\lib\site-packages\seaborn\_decorators.py:36:
```

```
FutureWarning: Pass the following variable as a keyword arg: x. From
version 0.12, the only valid positional argument will be `data`, and
passing other arguments without an explicit keyword will result in an
error or misinterpretation.
```


```
warnings.warn(
```

```
<AxesSubplot:xlabel='Gender', ylabel='count'>
```



```
plt.hist(data.Age)


(array([31., 19., 34., 29., 16., 26., 15., 10.,  6., 14.]),
 array([18. , 23.2, 28.4, 33.6, 38.8, 44. , 49.2, 54.4, 59.6, 64.8,
       70. ]),
 <BarContainer object of 10 artists>)
```



```
sns.scatterplot(data['Annual Income (k$)'],data['Spending Score (1-100)'])
```


```
D:\Anaconda\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variables as keyword args: x, y.
From version 0.12, the only valid positional argument will be `data`,
and passing other arguments without an explicit keyword will result in
an error or misinterpretation.
  warnings.warn(
```

```
<AxesSubplot:xlabel='Annual Income (k$)', ylabel='Spending Score (1-
100)'\>
```




```
plt.boxplot(data.iloc[0:,0])
```

```
{'whiskers': [<matplotlib.lines.Line2D at 0x1a310b28940>,
<matplotlib.lines.Line2D at 0x1a310b28c10>],
'caps': [<matplotlib.lines.Line2D at 0x1a310b28fa0>,
<matplotlib.lines.Line2D at 0x1a310b38370>],
'boxes': [<matplotlib.lines.Line2D at 0x1a310b284f0>],
'medians': [<matplotlib.lines.Line2D at 0x1a310b38700>],
'fliers': [<matplotlib.lines.Line2D at 0x1a310b38a90>],
'means': []}
```




```
plt.boxplot(data.iloc[0:,2])
```

```
{'whiskers': [<matplotlib.lines.Line2D at 0x1a310e8b2e0>,  
<matplotlib.lines.Line2D at 0x1a310e8b5e0>],  
'caps': [<matplotlib.lines.Line2D at 0x1a310e8b940>,  
<matplotlib.lines.Line2D at 0x1a310e8bcd0>],  
'boxes': [<matplotlib.lines.Line2D at 0x1a310e7df70>],  
'medians': [<matplotlib.lines.Line2D at 0x1a310e940a0>],  
'fliers': [<matplotlib.lines.Line2D at 0x1a310e94430>],  
'means': []}
```




```
plt.boxplot(data.iloc[0:,3])
```

```
{'whiskers': [<matplotlib.lines.Line2D at 0x1a310ee9be0>,  
             <matplotlib.lines.Line2D at 0x1a310ee9f70>],  
 'caps': [<matplotlib.lines.Line2D at 0x1a310ef6340>,  
          <matplotlib.lines.Line2D at 0x1a310ef66d0>],  
 'boxes': [<matplotlib.lines.Line2D at 0x1a310ee9910>],  
 'medians': [<matplotlib.lines.Line2D at 0x1a310ef6a60>],  
 'fliers': [<matplotlib.lines.Line2D at 0x1a310ef6df0>],  
 'means': []}
```



```
plt.boxplot(data.iloc[0:,4])
```

```
{'whiskers': [<matplotlib.lines.Line2D at 0x1a310f54a60>,  
             <matplotlib.lines.Line2D at 0x1a310f54df0>],  
 'caps': [<matplotlib.lines.Line2D at 0x1a310f631c0>,  
          <matplotlib.lines.Line2D at 0x1a310f63550>],  
 'boxes': [<matplotlib.lines.Line2D at 0x1a310f546d0>],  
 'medians': [<matplotlib.lines.Line2D at 0x1a310f638e0>],  
 'fliers': [<matplotlib.lines.Line2D at 0x1a310f63c70>],  
 'means': []}
```



```
data = pd.get_dummies(data,columns=["Gender"])
```

```
data
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)	\
0	1	19	15	39	
1	2	21	15	81	
2	3	20	16	6	
3	4	23	16	77	
4	5	31	17	40	
...	...	...	...	...	
195	196	35	120	79	
196	197	45	126	28	
197	198	32	126	74	
198	199	32	137	18	
199	200	30	137	83	

	Gender_Female	Gender_Male
0	0	1
1	0	1
2	1	0
3	1	0
4	1	0
...	...	...
195	1	0
196	1	0
197	0	1
198	0	1
199	0	1

```

[200 rows x 6 columns]
data.drop('CustomerID',axis=1,inplace=True)
x = data.drop('Spending Score (1-100)',axis=1)
y = data['Spending Score (1-100)']

trainx,testx,trainy,testy =
x.iloc[0:150,0:],x.iloc[150:,0:],y.iloc[0:150],y.iloc[150:]

model1 = RandomForestClassifier().fit(trainx,trainy)
model1.score(testx,testy)

0.04

model2 = GradientBoostingClassifier().fit(trainx,trainy)
model2.score(testx,testy)

0.02

model3 = AdaBoostClassifier().fit(trainx,trainy)
model3.score(testx,testy)

0.02

model3.predict(testx)

array([90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90,
90,
      90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90,
90,
      90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90, 90,
90],
      dtype=int64)

pred_y = model1.predict(testx)
mean_squared_error(testy,pred_y)

1749.56

```